


Neural Networks

October 4, 2017

Inspiration



Feed-forward Artificial Neural Network (ANN)



Motivating example: XOR (Exclusive Or)




$$\mathcal{O} = (A + B) \cdot \overline{(A \cdot B)}$$

A	B	\mathcal{O}
0	0	0
0	1	1
1	0	1
1	1	0

Choosing a linear model:

$$f(\mathbf{x}; \mathbf{w}, b) = \mathbf{x}^T \mathbf{w} + b$$

we get $\mathbf{w} = \mathbf{0}$ and $b = \frac{1}{2}$, i.e. the linear model outputs 0.5 everywhere.



Learning with simple Feedforward Network

- One hidden layer with two units.

- Two step process:



- ▶ $\mathbf{h} = f^{(1)}(\mathbf{x}; \mathbf{W}, \mathbf{c})$

- ▶ $y = f^{(2)}(\mathbf{h}; \mathbf{w}, b)$

- The complete model looks like:

$$f(\mathbf{x}; \mathbf{W}, \mathbf{c}, \mathbf{w}, b) = f^{(2)}(f^{(1)}(\mathbf{x}))$$

- What should $f^{(1)}$ and $f^{(2)}$ look like?



Choosing both $f^{(1)}$ and $f^{(2)}$ as linear models we end up with a linear model.

$$f^{(1)} = \mathbf{W}^T \mathbf{x}, \quad f^{(2)} = \mathbf{h}^T \mathbf{w} \quad \implies \quad f(\mathbf{x}) = \mathbf{w}^T \mathbf{W}^T \mathbf{x}$$

Clearly we need one of the functions to be non-linear!


Rectified Linear Unit (ReLU)

- We want to choose $f^{(1)}$ as a non-linear function.
- Modern neural networks employ Rectified linear units.

$$g(z) = \max(0, z)$$

- We then get:

$$f^{(1)} = \max(0, \mathbf{W}^T \mathbf{x} + \mathbf{c})$$



Our complete model becomes:

$$f(\mathbf{x}; \mathbf{W}, \mathbf{c}, \mathbf{w}, b) = \mathbf{w}^T \max(0, \mathbf{W}^T \mathbf{x} + \mathbf{c}) + b$$


The complete XOR model

$$f(\mathbf{x}; \mathbf{W}, \mathbf{c}, \mathbf{w}, b) = \mathbf{w}^T \max(0, \mathbf{W}^T \mathbf{x} + \mathbf{c}) + b$$

$$\mathbf{W} = \begin{bmatrix} 1 & 1 \\ 1 & 1 \end{bmatrix} \quad \mathbf{c} = \begin{bmatrix} 0 \\ -1 \end{bmatrix} \quad \mathbf{w} = \begin{bmatrix} 1 \\ -2 \end{bmatrix} \quad b = 0$$

$$X: \begin{bmatrix} 0 & 0 \\ 0 & 1 \\ 1 & 0 \\ 1 & 1 \end{bmatrix} \xrightarrow{\mathbf{W}} \begin{bmatrix} 0 & 0 \\ 1 & 1 \\ 1 & 1 \\ 2 & 2 \end{bmatrix} \xrightarrow{+\mathbf{c}} \begin{bmatrix} 0 & -1 \\ 1 & 0 \\ 1 & 0 \\ 2 & 1 \end{bmatrix} \xrightarrow{\text{ReLU}} \begin{bmatrix} 0 & 0 \\ 1 & 0 \\ 1 & 0 \\ 2 & 1 \end{bmatrix} \xrightarrow{\mathbf{w}^T} \begin{bmatrix} 0 \\ 1 \\ 1 \\ 0 \end{bmatrix}$$

General Feedforward network



- A feedforward network consist of layers of neurons.
- The output of neurons in layer n is passed as input to the neurons in layer $n + 1$

Neuron/Perceptron

- 1 **Perceptron** receives input

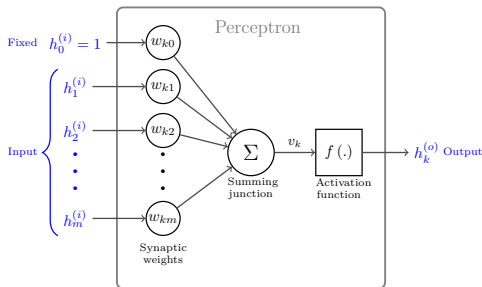
$\mathbf{h}^{(i)} = (h_1^{(i)}, h_2^{(i)}, \dots, h_m^{(i)})$ from previous layer. If a **bias** is included in the model add a fixed $h_0^{(i)} = 1$ to the input.

- 2 Inputs are multiplied by the **Synaptic Weights**, and added together at the **Summing Junction** (Linear/Affine transformation).



$$v_k = \mathbf{w}_k^T \mathbf{h}^{(i)} = \sum_{j=0}^m w_{kj} h_j^{(i)}$$

- 3 **Activation function** is evaluated to give **Perceptron** output (Non-linear transformation).

$$h_k^{(o)} = f(v_k)$$



Feeding forward



- Each Perceptron in layer n sends its output to each Perceptron in layer $n + 1$.
- Each Perceptron in layer $n + 1$ receives input from all Perceptrons in layer n .

Evaluating a feedforward network

Evaluating Perceptron k in layer n

$$h_k^{(n)} = f_k^{(n)} \left(\mathbf{w}_k^{(n)T} \mathbf{h}^{(n-1)} \right)$$

Evaluating the whole network is just a question of evaluating each layer in succession, i.e. **Feed-forward**.

$$h_k^{(1)} = f_k^{(1)} \left(\mathbf{w}_k^{(1)T} \mathbf{h}^{(0)} \right)$$

$$h_k^{(2)} = f_k^{(2)} \left(\mathbf{w}_k^{(2)T} \mathbf{h}^{(1)} \right)$$

...


$$h_k^{(O)} = f_k^{(O)} \left(\mathbf{w}_k^{(O)T} \mathbf{h}^{(O-1)} \right)$$

If all neurons in a layer has the same activation function we can rewrite the layer evaluation:


$$\mathbf{h}^{(n)} = f^{(n)} \left(\mathbf{W}^{(n)T} \mathbf{h}^{(n-1)} \right)$$

Activation functions


Rectified Linear Unit (ReLU)




Leaky/Parametric Rectified Linear Unit (PReLU)




Absolute Rectified Linear Unit




Sigmoid



Hyperbolic Tangent



Softplus



Minimization

Optimizing the weight of a feedforward neural network is a minimization problem!

Optimization

To optimize the weights of our network, we want, given a training set $x \in \mathbb{X}$, to minimize a loss function $J(w)$ over the training set. This could e.g. be a sum of squares loss function

$$J(w) = \frac{1}{2} \sum_{x \in \mathbb{X}} (f^*(x) - f(x; w))^2$$

The equation is at an optimum if the gradient is w.r.t. the weights is 0

$$\frac{\partial}{\partial w} J(w) = 0$$

We thus need derivatives of the loss function with respect to all linear weights w_{kj} of the Perceptrons.

$$\frac{\partial J(w)}{\partial w_{kj}} = 0$$

Back-propagation: Chain-rule

Chain-rule

Let $\mathbf{x} \in \mathbb{R}^m$, $\mathbf{y} \in \mathbb{R}^n$, $g : \mathbb{R}^m \rightarrow \mathbb{R}^n$, and $f : \mathbb{R}^n \rightarrow \mathbb{R}$. If $\mathbf{y} = g(\mathbf{x})$ and $z = f(\mathbf{y})$, then:


$$\frac{\partial z}{\partial x_i} = \sum_j \frac{\partial z}{\partial y_j} \frac{\partial y_j}{\partial x_i}$$

Written in vector notation we get,

$$\nabla_{\mathbf{x}} z = \left(\frac{\partial \mathbf{y}}{\partial \mathbf{x}} \right)^T \nabla_{\mathbf{y}} z,$$

where $\frac{\partial \mathbf{y}}{\partial \mathbf{x}}$ is the $n \times m$ Jacobian of g .

Graph derivative: Example



Derivative

$$\begin{aligned}\frac{\partial z}{\partial w} &= \frac{\partial z}{\partial y} \frac{\partial y}{\partial x} \frac{\partial x}{\partial w} \\ &= f'(y) f'(x) f'(w) \\ &= f'(f(f(w))) f'(f(w)) f'(w)\end{aligned}$$

Backpropagation derivation

Definition $\alpha_k^{(n)}$

Remember that the output of layer n can be written as

$$h_k^{(n)} = f_k^{(n)} \left(\mathbf{w}_k^{(n)T} \mathbf{h}^{(n-1)} \right) = f_k^{(n)} \left(\alpha_k^{(n)} \right)$$

with the definition

$$\alpha_k^{(n)} = \mathbf{w}_k^{(n)T} \mathbf{h}^{(n-1)} = \sum_i w_{ki}^{(n)} h_i^{(n-1)}$$

As $w_{ki}^{(n)}$ only enters the loss function through $\alpha_k^{(n)}$ we can, using the Chain rule, write the gradient of the loss function as:

$$\frac{\partial J(\mathbf{w})}{\partial w_{ki}^{(n)}} = \frac{\partial J(\mathbf{w})}{\partial \alpha_k^{(n)}} \frac{\partial \alpha_k^{(n)}}{\partial w_{ki}^{(n)}} = \delta_k^{(n)} \frac{\partial \alpha_k^{(n)}}{\partial w_{ki}^{(n)}}$$

where in the second step we have defined

$$\delta_k^{(n)} = \frac{\partial J(\mathbf{w})}{\partial \alpha_k^{(n)}}$$

Deriving and expression for $\delta_k^{(n)}$

First we apply the Chain rule:

$$\delta_k^{(n)} = \frac{\partial J(\mathbf{w})}{\partial \alpha_k^{(n)}} = \sum_l \frac{\partial J(\mathbf{w})}{\partial \alpha_l^{(n+1)}} \frac{\partial \alpha_l^{(n+1)}}{\partial \alpha_k^{(n)}}$$

First term is easy, it is just $\delta_l^{(n+1)}$. The second term gives:

$$\frac{\partial \alpha_l^{(n+1)}}{\partial \alpha_k^{(n)}} = \mathbf{w}_l^{(n)T} \frac{\partial \mathbf{h}^{(n)}}{\partial \alpha_k^{(n)}} = \sum_i w_{li}^{(n+1)} \frac{\partial h_i^{(n)}}{\partial \alpha_k^{(n)}} = w_{lk}^{(n+1)} h_k'^{(n)}$$

Combining first and second term we the expression:

$$\delta_k^{(n)} = h_k'^{(n)} \sum_l w_{lk}^{(n+1)} \delta_l^{(n+1)}$$

We see here why the method is known as **Back-propagation**, as $\delta_k^{(n)}$ depends on layer $n + 1$!

Backpropagated gradient

We now have an expression for the first term of our gradient expression:

$$\frac{\partial J(\mathbf{w})}{\partial w_{ki}^{(n)}} = \delta_k^{(n)} \frac{\partial \alpha_k^{(n)}}{\partial w_{ki}^{(n)}}$$

The second term is easy (remember that $\alpha_k^{(n)} = \sum_i w_{ki}^{(n)} h_i^{(n-1)}$):

$$\frac{\partial \alpha_k^{(n)}}{\partial w_{ki}^{(n)}} = h_i^{(n-1)}$$

Inserting this into our gradient expression we get an expression for our backpropagated gradient

$$\frac{\partial J(\mathbf{w})}{\partial w_{ki}^{(n)}} = \delta_k^{(n)} h_i^{(n-1)}$$

We can now use any gradient optimization algorithm (in fact we could also derive an expression for the Hessian. I leave this as an exercise ☺).

Universal Approximation Theorem

How general is the Neural Network model?

The **universal approximation theorem** (Hornik et al.) states that a feedforward network with a linear input layer and at least one hidden layer with *any* "squashing" activation function can approximate any Borel measurable function from one finite dimensional space to another with any desired non-zero amount of error, provided that the network is given enough hidden units.

$$f : \mathbb{R}^n \rightarrow \mathbb{R}^m$$



- Regardless of what function we are trying to learn, we know that a large neural network will be able to *represent* this.
- However, we are **not** guaranteed that the training algorithm will be able to *learn* this function.
- The theorem does not state how large the layer should be, and in **worst case** we need one hidden unit for each input configuration.

Architecture Design

Some architectures build a main chain, and then add extra features to it. These could include:

- Removing some connections between neurons to reduce number of parameters.
- Adding extra skip connections going from layer n to layer $n + 2$.
- Choosing which activation function to use. One could utilize different activation functions in each layer.
- Many problems need specialized neural network architectures, e.g. the Convolutional Neural Networks (CNNs) used in computer vision.

The depth of the network is also important. **Deeper** networks using many layers tend to perform better.

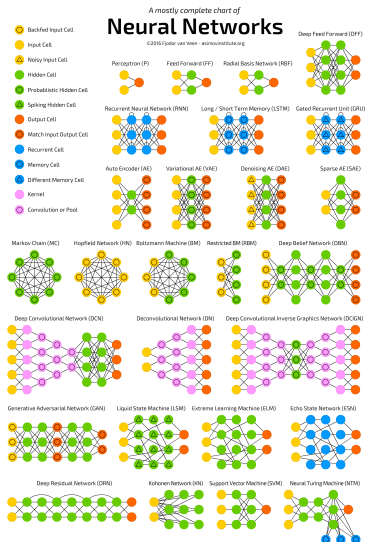


To sum up


- A neural network is a function that takes a set of input $\{x_i\}$, and return a set of output $\{y_i\}$.
- This is done through a series of Linear transformations using linear weights w , and non-linear function evaluations.
- Optimal weight parameters are found using an optimization/fitting algorithm, either using gradient or Hessian information.
- Network is evaluated by forward propagation.
- Derivatives are subsequently found by back propagation.
- Once optimal weight are found they can be stored and used for evaluating the network later.

Take home message: The Neural Network model is simply a non-linear function from a set of input variables $\{x_i\}$ to a set of output variables $\{y_i\}$ controlled by a vector w of adjustable parameters.

Flavors of Neural Networks



Resources



Available online: <http://www.deeplearningbook.org/>

And in pdf: <https://github.com/janishar/mit-deep-learning-book-pdf>

And with that...

... we are done for today!